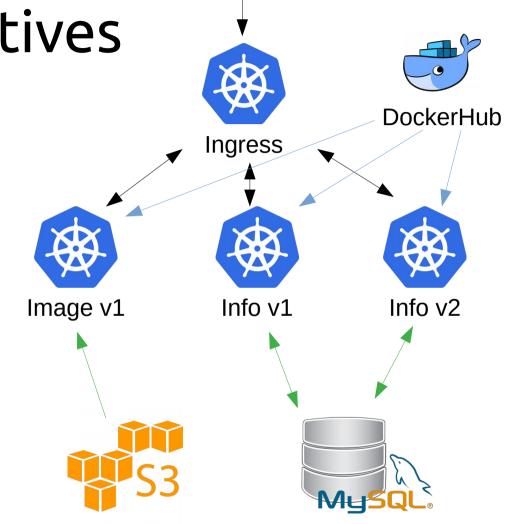
CLOUD COMPUTING PROJECT 1 - Part II

Lorenzo Leonini lorenzo.leonini@unine.ch

Objectives

- Deploy an API in the cloud
 - Using kubernetes!
- Leverage PaaS
 - Docker Images repository
 - SQL DB
 - NoSQL DB (S3)
 - Container orchestration
- Automate the process
 - Code => Deployment



Microservice - Watch info v2

- Compared to v1, minor changes + new search semantic
 - See README in git
 - New microservice
 - Duplicate v1 and apply the changes
 - Or decide between v1 and v2 at launch via env parameter
- Keep v1
 - Backward compatibility

Microservice - Image info v1

- New microservice
- Very simple (image proxy)
 - One endpoint /image/v1/get/{sku}
 - Will transfert an image from a NoSQL repo (AWS S3) to the user
 - Fixed mime type (image/png)
 - No image processing (ATM ;-)
 - Set proper caching headers and ETag
 - Lifetime 1h

Watches images

- The images are stored in a AWS S3 bucket (localized in Ireland)
 - Publicly accessible
 - Image format is transparent PNG (max 1024x1024)
 - https://s3-eu-west-1.amazonaws.com/cloudcomputing-2018/project1/ images/<SKU>.png
 - https://s3-eu-west-1.amazonaws.com/cloudcomputing-2018/project1/ images/CAC1111.BA0850.png
 - You can access them from GKE
 - Not in the same datacenter will add a slight delay
 - Not all images will be available! (~ 50% of SKUs have images)
 - Check for 200 or 404 HTTP codes

Git repository

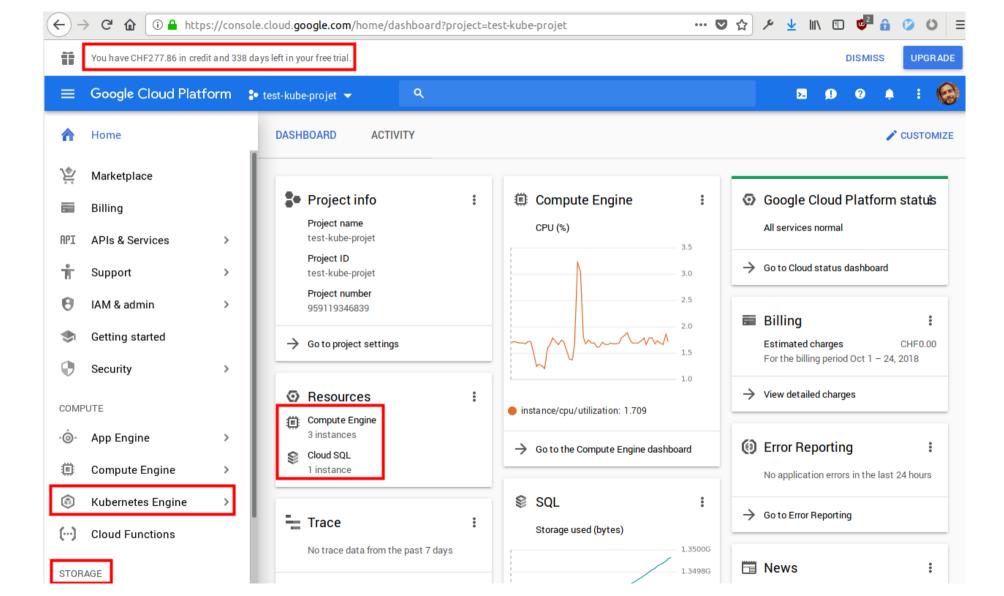
- Renaming
 - openapi.yaml to info_openapi_v1.yaml
 - swagger.yaml to info_swagger_v1.yaml
- Added
 - info_openapi_v2.yaml
 - image_openapi_v1.yaml
 - There is not anymore a swagger version
 - README to summarize changes in the repo

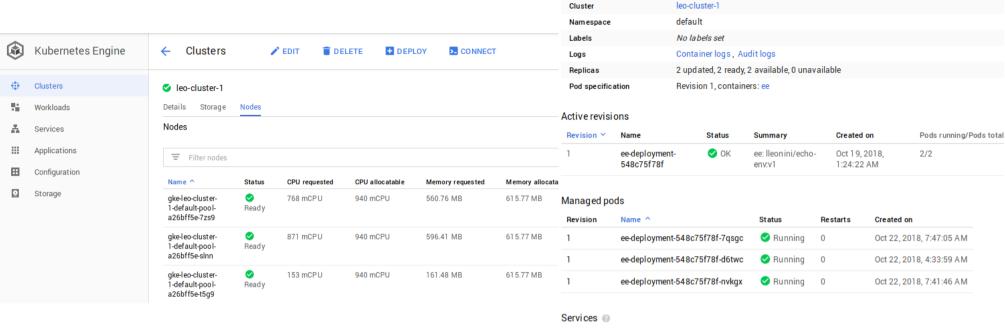
DockerHub

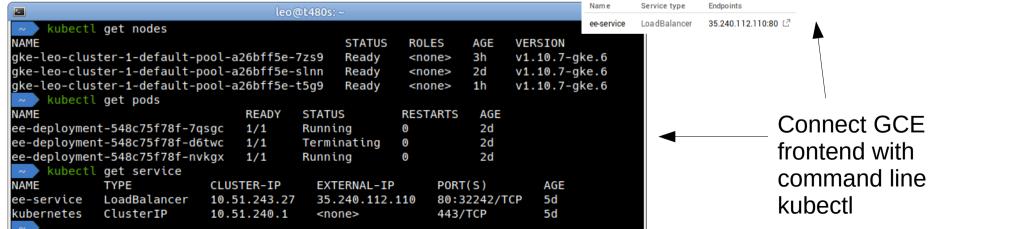
- Docker/Kubernetes/AWS can use any Docker images repositories
 - But DockerHub is the default
 - If you publish your images publicly on DockerHub, you will not have to manage complex access to a particular repo (and no authentication to manage)
 - Create a free account
- Microservices images to publish
 - Infos v1
 - Infos v2
 - Images v1
- Publishing updated images should be part of your build process

Kubernetes - GKE

- Create a Google Compute/Kubernetes Engine account
 - https://cloud.google.com/compute/
 - 300\$ / 1 year free credits
 - Create a kubernetes cluster in Europe West (close to AWS S3 for images)
 - Link your account to kubectl
 - https://cloud.google.com/kubernetes-engine/docs/quickstart
 - \$ gcloud container clusters get-credentials <cluster> --zone europe-west1-b --project
 <project>
 - Then, use kubectl as you will use it locally with minikube!
- On AWS, kubernetes also exists but is expensive (~200\$ / month with mininal config) and much complicated to deploy
 - Kubernetes is a technology from Google and it is much more integrated in GCE







SQL Database

- The SQL database must be shared between info services v1 and v2
- 2 possibilities (you can decide)
 - 1. Use a cloud managed MySQL (in GCE or in AWS)
 - Open firewall to access it from kubernetes pods
 - 2. Use kubernetes with MySQL image and a persistent volume
 - https://kubernetes.io/docs/tasks/run-application/run-single-instance-stateful-application/
 - Can be more difficult to initialize the DB
- Bonus
 - (re-) Initialize the DB automatically with deployments

Kubernetes - Ingress controller

- Use ingress controller to route to your microservices
 - /watch/v1/*
 - /watch/v2/*
 - /image/v1/*
- Ask for a static IP address
 - https://cloud.google.com/kubernetes-engine/docs/ tutorials/http-balancer

Kubernetes - Deployment

- Create 3 services (info v1, v2 and image)
 - => 3 deployment + replicaSets
 - Permits scalling instances
 - Permits rolling upgrade
- Create Ingress controller to route to these 3 services
- Bonus
 - Container probes
 - Add a liveness check (HTTP query returning 200)
 - https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/

Deliverables - Delay - Grading

- Deliverables
 - IP/DNS to test your API!
 - Short documentation (README)
 - How to (re-)deploy
 - Script(s) to build & deploy
 - => rebuild 3 service images
 - => update 3 images to DockerHub
 - => kubernetes rolling upgrade
 - Kubernetes deployment file(s) yaml
 - Deployment instructions
 - Multiple object descriptions can be grouped in one file, use '\n---\n' as separator

- Delay: 3 weeks
- Grading
 - Try to automate the process as much as possible
 - Code changes to "in production" must be as simple as possible
 - Bonus: each + 0.5 points